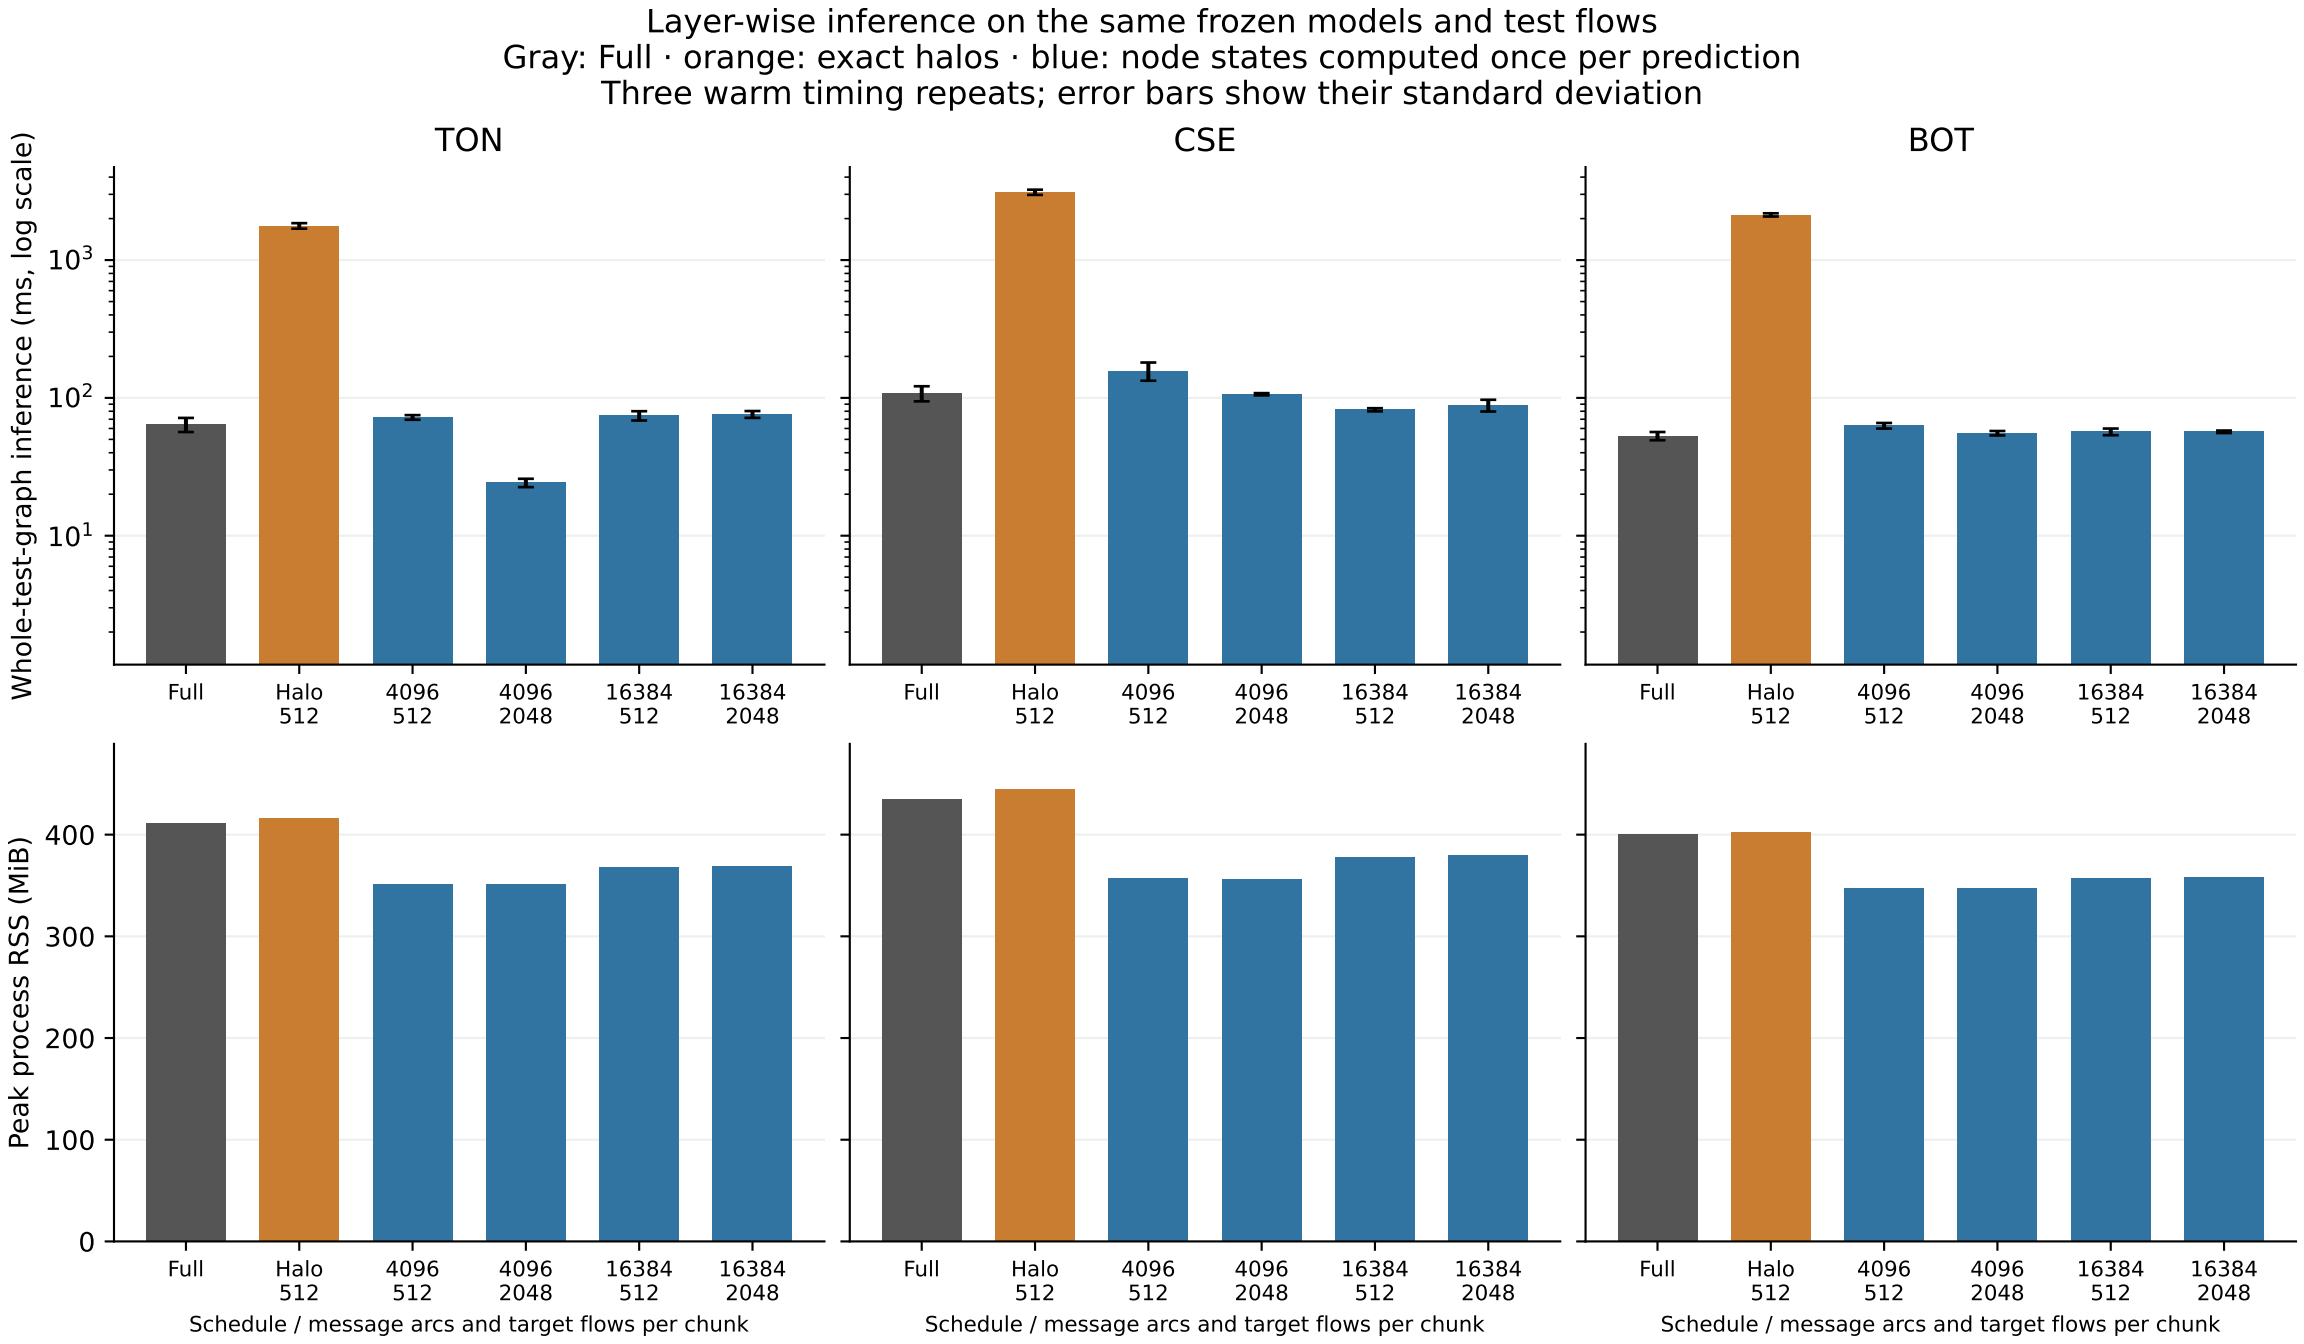
Layer-wise inference on the same frozen models and test flows
Gray: Full · orange: exact halos · blue: node states computed once per prediction
Three warm timing repeats; error bars show their standard deviation
TON
CSE
BOT
Whole-test-graph inference (ms, log scale)
10^3
10^2
10^1
Peak process RSS (MiB)
400
300
200
100
0
Full
Halo 512
4096 512
4096 2048
16384 512
16384 2048
Schedule / message arcs and target flows per chunk