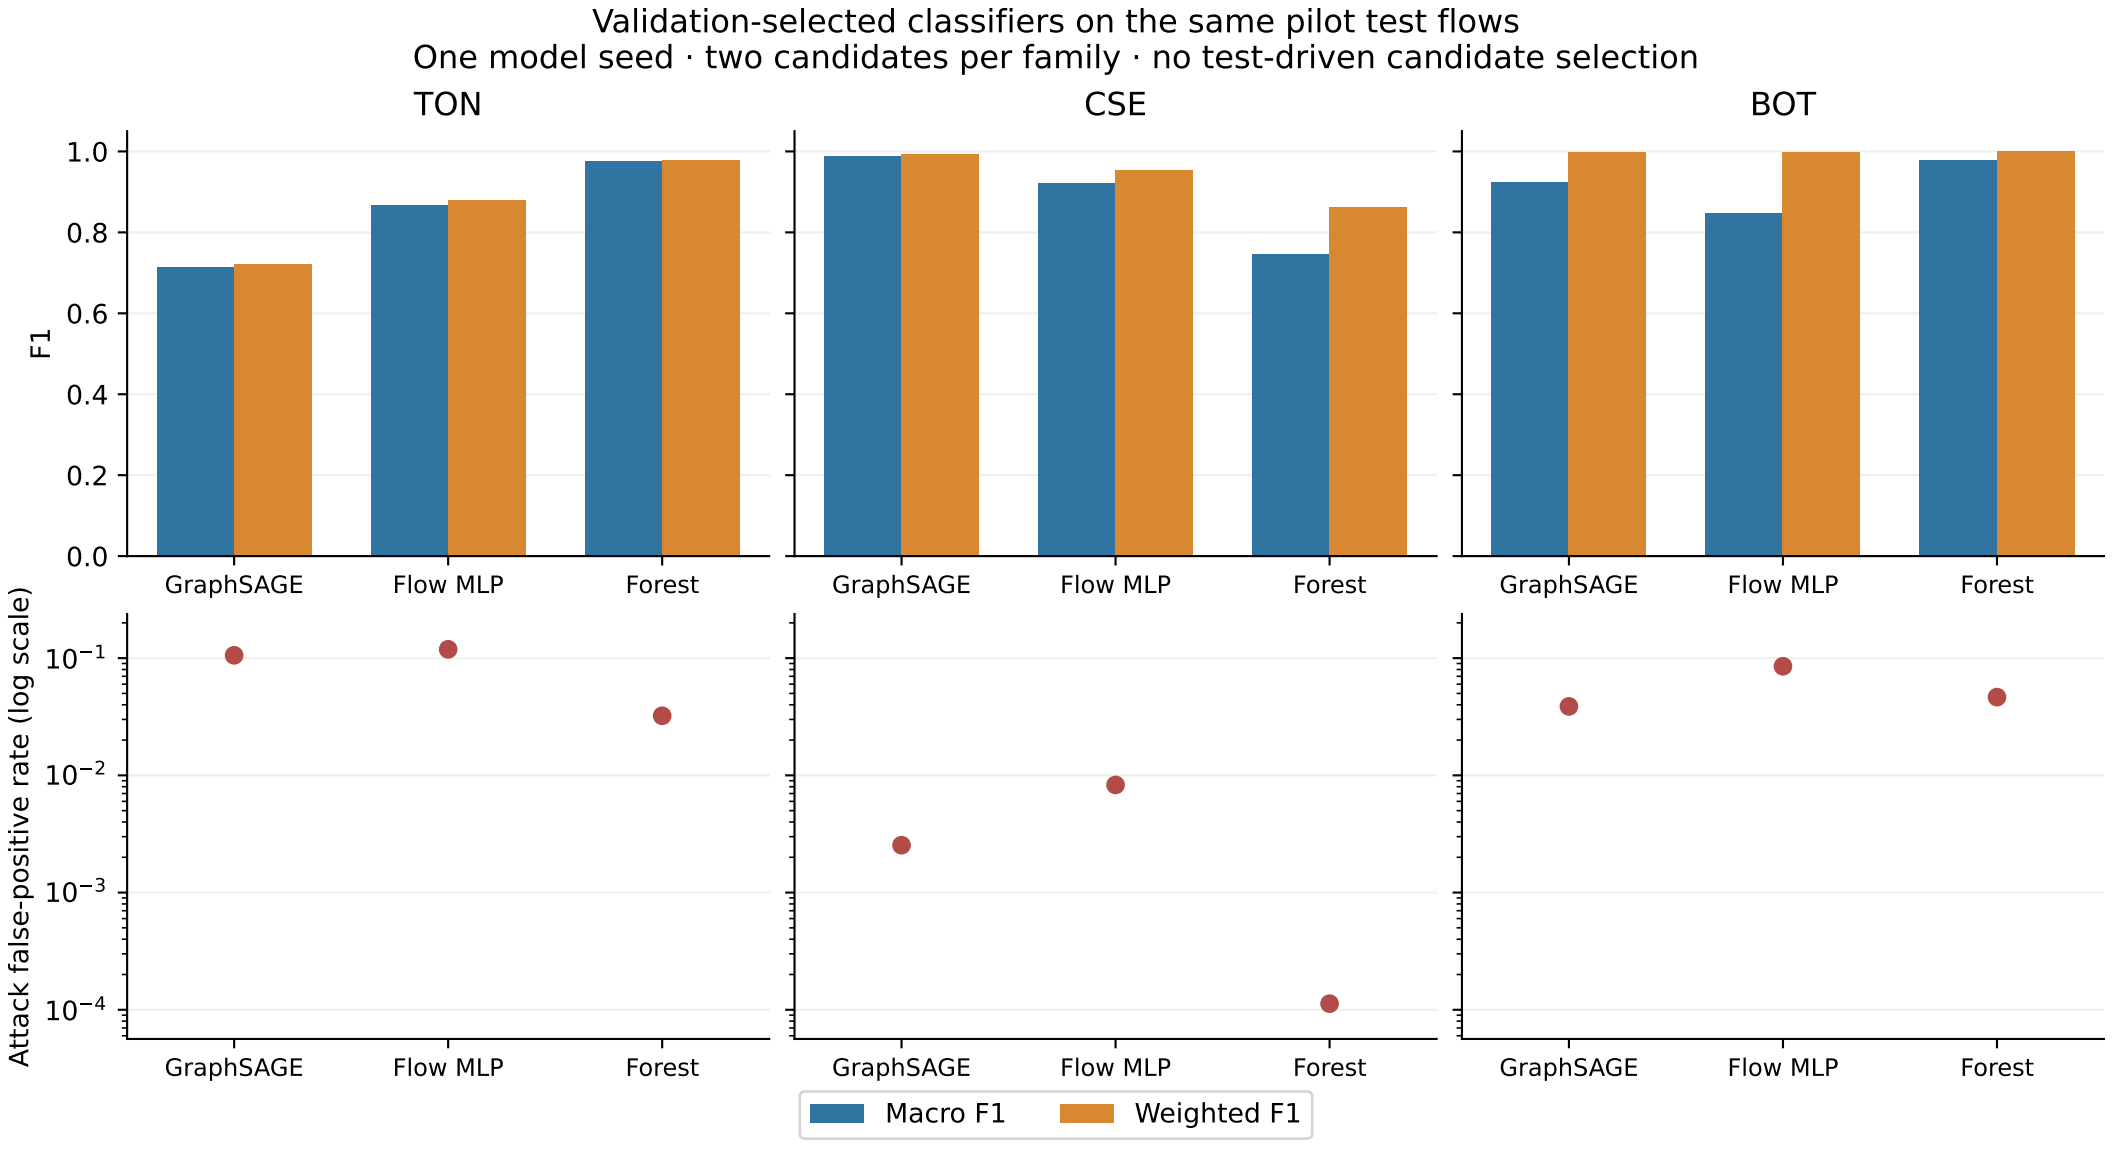
Validation-selected classifiers on the same pilot test flows
One model seed · two candidates per family · no test-driven candidate selection
TON
CSE
BOT
F1
1.0
0.8
0.6
0.4
0.2
0.0
GraphSAGE
Flow MLP
Forest
Attack false-positive rate (log scale)
$10^{-1}$
$10^{-2}$
$10^{-3}$
$10^{-4}$
Macro F1
Weighted F1